

HW2

112062706 林泳成

- [HW2](#)
 - [Implementation](#)
 - [pthread](#)
 - [hybrid](#)
 - [Experiment & Analysis](#)
 - [Performance matrix & Speedup](#)
 - [Profile](#)
 - [Discussion](#)
 - [Conclusion](#)

Implementation

pthread

```
int curX = 0, curY = 0, blockPerThread = 2500;

int jobRemain(){
    // Return available block size
    // Guarantee cur not pass limit
    int totalBlock = width * height;
    int curSpanX = curX + curY * width;
    if(curSpanX + blockPerThread < totalBlock)
        return blockPerThread;
    else if(curSpanX >= totalBlock)
        return -1;
    else
        return totalBlock - curSpanX;
}
```

```

void* calculate(void*){
    while(true){
        pthread_mutex_lock(&mutex);
        int block = jobRemain();
        int startX, startY;
        if(block == -1){
            pthread_mutex_unlock(&mutex);
            break;
        }
        else{
            startX = curX;
            startY = curY;
            curX += block;
            while(curX >= width){
                curX -= width;
                ++curY;
            }
            pthread_mutex_unlock(&mutex);
        }
    }
}

```

每個thread都會在while迴圈內去拿新的工作，工作大小定義在global variable(blockPerThread)。每個thread都會一直拿工作直到沒有新的工作可以拿。jobRemain function會回傳可以拿的資料量。另外有用curX和curY指向下一個工作的起始位置，所以會用lock把這個curX和curY的更新包起來。

```

if(block != blockPerThread){
    int iter = 0;
    while(true){
        if(iter == block)
            break;
        int repeats = 0;
        double x = 0;
        double y = 0;
        double length_squared = 0;
        double y0 = startY * ((upper - lower) / height) + lower;
        double x0 = startX * ((right - left) / width) + left;
        while (repeats < iters && length_squared < 4) {
            double temp = x * x - y * y + x0;
            y = 2 * x * y + y0;
            x = temp;
            length_squared = x * x + y * y;
            ++repeats;
        }
        image[startY * width + startX] = repeats;
        ++startX;
        if(startX == width){
            ++startY;
            startX = 0;
        }
        ++iter;
    }
    break;
}

```

如果拿到的工作是最後一份，那就讓他直接執行
 剩下的部分會用vectorization實作

```

double startXVec[8], startYVec[8], repeatCpy[8], startYCpy[8], startXCpy[8], xCpy[8], yCpy[8], lengthSquaredCpy[8], x0Cpy[8], y0Cpy[8];
for(int i=0;i<8;++i){
    // ++iter;
    startXVec[i] = (double)startX;
    startYVec[i] = (double)startY;
    ++startX;
    if(startX >= width){
        startX = 0;
        ++startY;
    }
}

__m512d mxVec = _mm512_set1_pd(0.0);
__m512d myVec = _mm512_set1_pd(0.0);
__m512d mstartXVec = _mm512_loadu_pd(startXVec);
__m512d mstartYVec = _mm512_loadu_pd(startYVec);
__m512d mlengthSquared = _mm512_set1_pd(0.0);
__m512d mupperVec = _mm512_set1_pd(upper);
__m512d mlowerVec = _mm512_set1_pd(lower);
__m512d mrightVec = _mm512_set1_pd(right);
__m512d mleftVec = _mm512_set1_pd(left);
__m512d mheightVec = _mm512_set1_pd((double)height);
__m512d mwidthVec = _mm512_set1_pd((double)width);

__m512d mitersVec = _mm512_set1_pd((double)iters);
__m512d mrepeatsVec = _mm512_set1_pd(0.0);

__m512d mx0Vec = _mm512_fmadd_pd(mstartXVec, _mm512_div_pd(_mm512_sub_pd(mrightVec, mleftVec), mwidthVec), mleftVec);
__m512d my0Vec = _mm512_fmadd_pd(mstartYVec, _mm512_div_pd(_mm512_sub_pd(mupperVec, mlowerVec), mheightVec), mlowerVec);

```

先把前8個資料塞進__m512d的vector，之後就同時更新8個變數

```

while(iter < block){
    __mmask8 mask4 = __mm512_cmp_pd_mask(mlengthSquared, __mm512_set1_pd(4.0), _CMP_GE_OQ);
    __mmask8 maskRepeats = __mm512_cmp_pd_mask(mrepeatsVec, mitersVec, _CMP_GE_OQ);
    __mmask8 maskEmpty = __mm512_cmp_pd_mask(mrepeatsVec, __mm512_set1_pd(-1.0), _CMP_EQ_OQ);
    for(idx=0;idx<8;++idx){
        if((mask4 & (1 << idx)) || (maskRepeats & (1 << idx)) && !(maskEmpty & (1 << idx))){
            __mm512_store_pd(&xCpy, mxVec);
            __mm512_store_pd(&yCpy, myVec);
            __mm512_store_pd(&x0Cpy, mx0Vec);
            __mm512_store_pd(&y0Cpy, my0Vec);
            __mm512_store_pd(&repeatCpy, mrepeatsVec);
            __mm512_store_pd(&startXCpy, mstartXVec);
            __mm512_store_pd(&startYCpy, mstartYVec);
            image[(int)startYCpy[idx] * width + (int)startXCpy[idx]] = (int)repeatCpy[idx];
            if(iter == block)
                break;
            else{
                ++iter;
                __mm512_store_pd(&xCpy, mxVec);
                __mm512_store_pd(&yCpy, myVec);
                __mm512_store_pd(&x0Cpy, mx0Vec);
                __mm512_store_pd(&y0Cpy, my0Vec);
                __mm512_store_pd(&lengthSquaredCpy, mlengthSquared);
                repeatCpy[idx] = 0.0;
                startXCpy[idx] = (double)startX;
                startYCpy[idx] = (double)startY;
                xCpy[idx] = 0.0;
                yCpy[idx] = 0.0;
                lengthSquaredCpy[idx] = 0.0;
                x0Cpy[idx] = startX * ((right - left) / width) + left;
                y0Cpy[idx] = startY * ((upper - lower) / height) + lower;
                mrepeatsVec = __mm512_load_pd(&repeatCpy);
                mstartXVec = __mm512_load_pd(&startXCpy);
                mstartYVec = __mm512_load_pd(&startYCpy);
                mxVec = __mm512_load_pd(&xCpy);
                myVec = __mm512_load_pd(&yCpy);
                mlengthSquared = __mm512_load_pd(&lengthSquaredCpy);
                mx0Vec = __mm512_load_pd(&x0Cpy);
                my0Vec = __mm512_load_pd(&y0Cpy);
                ++startX;

```

```

                ++startX;
                if(startX >= width){
                    startX = 0;
                    ++startY;
                }
            }
        }
    }

    __m512d tmpVec = __mm512_add_pd(__mm512_sub_pd(__mm512_mul_pd(mxVec, mxVec), __mm512_mul_pd(myVec, myVec)), mx0Vec);
    myVec = __mm512_add_pd(__mm512_mul_pd(__mm512_set1_pd(2.0), __mm512_mul_pd(mxVec, myVec)), my0Vec);
    mxVec = tmpVec;
    mlengthSquared = __mm512_add_pd(__mm512_mul_pd(mxVec, mxVec), __mm512_mul_pd(myVec, myVec));
    mrepeatsVec = __mm512_add_pd(mrepeatsVec, __mm512_set1_pd(1.0));
}

```

每次都去檢查看看有沒有任務跑完了，如果有就把那一格的repeats次數寫出來，並且把新的工作覆蓋在空出來的格子上

```

for(int i=0;i<8;++i){
    if(i == idx)
        continue;
    int repeats = 0;
    double x = 0;
    double y = 0;
    double length_squared = 0;
    double y0 = (int)startYCpy[i] * ((upper - lower) / height) + lower;
    double x0 = (int)startXCpy[i] * ((right - left) / width) + left;
    while (repeats < iters && length_squared < 4) {
        double temp = x * x - y * y + x0;
        y = 2 * x * y + y0;
        x = temp;
        length_squared = x * x + y * y;
        ++repeats;
    }
    image[(int)startYCpy[i] * width + (int)startXCpy[i]] = repeats;
}

```

直到剩下最後7個任務無法繼續用vectorization，就重新sequential得跑

優化後

```

__mmask8 maskEmpty = _mm512_cmp_pd_mask(mrepeatsVec, _mm512_set1_pd(-1.0), _CMP_EQ_0Q);
for(idx=0;idx<8;++idx){
    if(((mask4 & (1 << idx)) || (maskRepeats & (1 << idx))) && !(maskEmpty & (1 << idx))){

```

其實後來發現vectorization不一定要全滿才能做，裡面也可以塞無用的空直，這樣就不需要重跑最後7個任務了(這個優化因為是每次拿新工作就會在最後執行到，所以跑judge的時候總共快了16秒左右)

```

for(int i=0;i<min(8, block);++i){
    // ++iter;
    startXVec[i] = (double)startX;
    startYVec[i] = (double)startY;
    ++startX;
    if(startX >= width){
        startX = 0;
        ++startY;
    }
}

```

```

if(block != blockPerThread)
    break;

```

還有一個小優化是在最後拿不滿blockPerThread的工作量時才會做的sequential run。如果也把這些塞不滿8格的工作也拿去做vectorization就可以不需要一個一個慢慢做(這個優化只優化最後一個工作，所以進步幅度有限，最後跑judge總共也只快了0.3秒)

```

// __mmask8 maskEmpty = _mm512_cmp_pd_mask(mrepeatsVec, _mm512_set1_pd(-1.0), _CMP_EQ_0Q);
for(idx=0;idx<8;++idx){
    if(((mask4 & (1 << idx)) || (maskRepeats & (1 << idx))){

```

後來發現如果圖片足夠大，那最後7個點的repeats次數一定會比重新計算(curX=0, curY+=1)還要早

算完，因此就不需要另外把不合法的工作排除，也就不需要開mask來遮掉有問題的工作(跑judge快了2秒)

hybrid

```
int floorHeight = height / size;
int extraRows = height % size;
int rows = rank < extraRows ? floorHeight + 1 : floorHeight;
```

每個process都會分到rows行的工作。如果無法整除，前面幾個process會拿到多一行的工作

```
#pragma omp parallel num_threads(CPU_COUNT(&cpu_set))
{
    int startX = 0, startY = 0;
    while(startY < rows){
        int idx = -1;
        int iter = 0;

        double startXVec[8], startYVec[8], repeatCpy[8], startYCpy[8], startXCpy[8], xCpy[8], yCpy[8], lengthSquaredCpy[8], x0Cpy[8], y0Cpy[8];
        for(int i=0;i<8;++i){
            ++iter;
            startXVec[i] = (double)startX;
            startYVec[i] = (double)startY;
            ++startX;
        }
        __m512d mxVec = _mm512_set1_pd(0.0);
        __m512d myVec = _mm512_set1_pd(0.0);
        __m512d mstartXVec = _mm512_loadu_pd(startXVec);
        __m512d mstartYVec = _mm512_loadu_pd(startYVec);
        __m512d mlengthSquared = _mm512_set1_pd(0.0);
        __m512d mupperVec = _mm512_set1_pd(upper);
        __m512d mlowerVec = _mm512_set1_pd(lower);
        __m512d mrighVec = _mm512_set1_pd(right);
        __m512d mleftVec = _mm512_set1_pd(left);
        __m512d mheightVec = _mm512_set1_pd((double)height);
        __m512d mwidthVec = _mm512_set1_pd((double)width);

        __m512d mitersVec = _mm512_set1_pd((double)iters);
        __m512d mrepeatsVec = _mm512_set1_pd(0.0);

        __m512d mx0Vec = _mm512_fmadd_pd(mstartXVec, _mm512_div_pd(_mm512_sub_pd(mrighVec, mleftVec), mwidthVec), mleftVec);
        __m512d my0Vec = _mm512_fmadd_pd(_mm512_add_pd(_mm512_mul_pd(mstartYVec, _mm512_set1_pd((double)size)), _mm512_set1_pd((double)rank)), _mm512_div_pd(

        while(iter < width){
            __mmask8 mask4 = _mm512_cmp_pd_mask(mlengthSquared, _mm512_set1_pd(4.0), _CMP_GE_QQ);
            __mmask8 maskRepeats = _mm512_cmp_pd_mask(mrepeatsVec, mitersVec, _CMP_GE_QQ);
            for(idx=0;idx<8;++idx){
                if((mask4 & (1 << idx)) || (maskRepeats & (1 << idx))){
                    _mm512_store_pd(&xCpy, mxVec);
                    _mm512_store_pd(&yCpy, myVec);
                    _mm512_store_pd(&x0Cpy, mx0Vec);
                    _mm512_store_pd(&y0Cpy, my0Vec);
                    _mm512_store_pd(&repeatCpy, mrepeatsVec);
```

OpenMP開出來的thread執行的工作基本和pthread一樣，都用vectorization去優化計算
每個process實際上會是輪流拿每一行的工作


```

        _mm512_store_pd(&repeatCpy, mrepeatsVec);
        _mm512_store_pd(&startXCpy, mstartXVec);
        _mm512_store_pd(&startYCpy, mstartYVec);
        image[(int)startYCpy[idx] * width + (int)startXCpy[idx]] = (int)repeatCpy[idx];
        if(iter == width)
            break;
        else{
            ++iter;
            _mm512_store_pd(&xCpy, mxVec);
            _mm512_store_pd(&yCpy, myVec);
            _mm512_store_pd(&x0Cpy, mx0Vec);
            _mm512_store_pd(&y0Cpy, my0Vec);
            _mm512_store_pd(&lengthSquaredCpy, mlengthSquared);
            repeatCpy[idx] = 0.0;
            startXCpy[idx] = (double)startX;
            startYCpy[idx] = (double)startY;
            xCpy[idx] = 0.0;
            yCpy[idx] = 0.0;
            lengthSquaredCpy[idx] = 0.0;
            x0Cpy[idx] = startX * ((right - left) / width) + left;
            y0Cpy[idx] = ((double)startY*size) + rank * ((upper - lower) / height) + lower;
            mrepeatsVec = _mm512_load_pd(&repeatCpy);
            mstartXVec = _mm512_load_pd(&startXCpy);
            mstartYVec = _mm512_load_pd(&startYCpy);
            mxVec = _mm512_load_pd(&xCpy);
            myVec = _mm512_load_pd(&yCpy);
            mlengthSquared = _mm512_load_pd(&lengthSquaredCpy);
            mx0Vec = _mm512_load_pd(&x0Cpy);
            my0Vec = _mm512_load_pd(&y0Cpy);
            ++startX;
        }
    }
}

__m512d tmpVec = _mm512_add_pd(_mm512_sub_pd(_mm512_mul_pd(mxVec, mxVec), _mm512_mul_pd(myVec, myVec)), mx0Vec);
myVec = _mm512_add_pd(_mm512_mul_pd(_mm512_set1_pd(2.0), _mm512_mul_pd(mxVec, myVec)), my0Vec);
mxVec = tmpVec;
mlengthSquared = _mm512_add_pd(_mm512_mul_pd(mxVec, mxVec), _mm512_mul_pd(myVec, myVec));
mrepeatsVec = _mm512_add_pd(mrepeatsVec, _mm512_set1_pd(1.0));
}

for(int i=0;i<8;++i){
    if(i == idx)
        continue;
    int repeats = 0;
    double x = 0;
    double y = 0;
    double length_squared = 0;
    double y0 = ((int)startYCpy[i]*size+rank) * ((upper - lower) / height) + lower;
    double x0 = (int)startXCpy[i] * ((right - left) / width) + left;
    while (repeats < iters && length_squared < 4) {
        double temp = x * x - y * y + x0;
        y = 2 * x * y + y0;
        x = temp;
        length_squared = x * x + y * y;
        ++repeats;
    }
    image[(int)startYCpy[i] * width + (int)startXCpy[i]] = repeats;
}
++startY;
startX = 0;

```

剩餘無法塞進__m512d的工作就在最後用sequential去做

```

displs[0] = 0;
for(int i = 0; i < size; i++){
    if(i < extraRows)
        recvcunts[i] = (floorHeight + 1) * width;
    else
        recvcunts[i] = floorHeight * width;

    if(i != 0)
        displs[i] = displs[i-1] + recvcunts[i-1];
}

```

```

MPI_Gatherv(image, rows*width, MPI_INT, gatherImage, recvcunts, displs, MPI_INT, 0, MPI_COMM_WORLD);

```

由於每個process是跳著拿工作的，且工作量都不一樣，所以直接用MPI_Gather不是個好方法
因此使用MPI_Gatherv來整合整張圖

MPI_Gatherv會需要多吃兩個參數，一個是用來存每個rank要傳的data量，一個是buffer接收每個rank的資料時的起始位置

```

if(rank==0){
    int* final_image = (int*)malloc(width * height * sizeof(int));
    clock_gettime(CLOCK_MONOTONIC, &start_t);
    int index = 0;
    for(int i = 0; i < rows; i++){
        int stride = rows;
        int extraRowLeft = extraRows;
        for(int row = i; row < height && index < height*width; row += stride){
            for(int col = 0; col < width; col++){
                final_image[index] = gatherImage[row*width + col];
                index++;
            }
            --extraRowLeft;
            if(extraRowLeft < 0){
                stride = floorHeight;
            }
        }
    }
}

```

Gather完的圖片會是每個process跳著拿工作的結果
所以需要再跳著把圖片組回來

優化後

```

#pragma omp parallel num_threads(CPU_COUNT(&cpu_set))
{
    double startXVec[8], startYVec[8], repeatCpy[8], startYCpy[8], startXcpy[8], xCpy[8], yCpy[8], lengthSquaredCpy[8], x0Cpy[8], y0Cpy[8];
    #pragma omp for schedule(static, 10)
    for(int startY = 0; startY < rows; ++startY){
        int startX = 0;
        for(int i=0; i<8; ++i){
            startXVec[i] = (double)startX;
            startYVec[i] = (double)startY;
            ++startX;
        }
        __m512d mxVec = _mm512_set1_pd(0.0);
        __m512d myVec = _mm512_set1_pd(0.0);
        __m512d mstartXVec = _mm512_loadu_pd(startXVec);
        __m512d mstartYVec = _mm512_loadu_pd(startYVec);
        __m512d mlengthSquared = _mm512_set1_pd(0.0);
        __m512d mupperVec = _mm512_set1_pd(upper);
        __m512d mlowerVec = _mm512_set1_pd(lower);
        __m512d mrightVec = _mm512_set1_pd(right);
        __m512d mleftVec = _mm512_set1_pd(left);
        __m512d mheightVec = _mm512_set1_pd((double)height);
        __m512d mwidthVec = _mm512_set1_pd((double)width);

        __m512d mitersVec = _mm512_set1_pd((double)iters);
        __m512d mrepeatsVec = _mm512_set1_pd(0.0);

        __m512d mx0Vec = _mm512_fmadd_pd(mstartXVec, _mm512_div_pd(_mm512_sub_pd(mrightVec, mleftVec), mwidthVec), mleftVec);
        __m512d my0Vec = _mm512_fmadd_pd(_mm512_add_pd(_mm512_mul_pd(mstartYVec, _mm512_set1_pd((double)size)), _mm512_set1_pd((double)rank)), _mm512_div_p
        int idx = -1;
        for(int iter = 8; iter < width; ++iter){
            while(true){
                __mmask8 mask4 = _mm512_cmp_pd_mask(mlengthSquared, _mm512_set1_pd(4.0), _CMP_GE_OQ);
                __mmask8 maskRepeats = _mm512_cmp_pd_mask(mrepeatsVec, mitersVec, _CMP_GE_OQ);
                bool breakFlag = false;
                for(idx=0; idx<8; ++idx){
                    if((mask4 & (1 << idx)) || (maskRepeats & (1 << idx))){
                        __mm512_store_pd(&xCpy, mxVec);
                        __mm512_store_pd(&yCpy, myVec);
                        __mm512_store_pd(&x0Cpy, mx0Vec);
                        __mm512_store_pd(&y0Cpy, my0Vec);
                        __mm512_store_pd(&repeatCpy, mrepeatsVec);
                        __mm512_store_pd(&startXCpy, mstartXVec);
                        __mm512_store_pd(&startYCpy, mstartYVec);
                    }
                }
            }
        }
    }
}

```

因為OpenMP的data parallelism需要把程式改成互相independent的for迴圈，所以我把原本的程式改成能用for迴圈來增加startY，讓原本會有dependency的程式能分離，讓OpenMP自動做工作分配(我原本跑judge要303秒，這個優化讓這隻程式變成跑judge只需要127秒。**突然發現OpenMP很吃programmer的使用**)

Experiment & Analysis

Performance matrix & Speedup

pthread

因為看到strict28到strict36都是一樣的iteration、圖寬和圖高，就只差在x0, x1, y0, y1，就想說拿來當作實驗input。結果發現時間差很多，最快的6秒就完成了，最慢的要107秒。

Input file	strict28.txt	strict29.txt	strict30.txt	strict31.txt	strict32.txt	strict33.txt	strict34.txt	strict35.txt	strict36.txt
x0	-0.6743255349	0.2936331265	-0.3421054598	-0.3070888275	-0.5506211524	-0.2989925066	-0.5506164692	-0.2931209325	-0.2872724083
x1	-0.674237311	0.2938268309	-0.2373971479	-0.247606819	-0.5506132349	-0.2772002993	-0.5506164628	-0.274142734	-0.2791228113
y0	0.3623627742	-0.01494847713	-0.6373595233	-0.6245844142	0.6273469513	-0.6327591639	0.6273445437	-0.6337125743	-0.6345413373
y1	0.362305991	-0.01505119437	-0.6884105744	-0.6535851592	0.6273427528	-0.6433840615	0.6273445404	-0.6429654881	-0.6385148108
Process	1	1	1	1	1	1	1	1	1
Thread	1	1	1	1	1	1	1	1	1
Initialization time	0.000019	0.00002	0.000018	0.00002	0.000022	0.00002	0.000019	0.000017	0.00002
Thread compute time	7.023348	24.733721	6.139337	69.2553	20.96561	83.905639	106.299448	46.753523	20.139368
Write PNG time	1.155882	0.947597	0.762026	0.968045	1.48397	0.900907	1.541206	1.000165	0.95577
Total time	8.179249	25.681338	6.901381	70.223365	22.449602	84.806566	107.840673	47.753705	21.095158
Process	1	1	1	1	1	1	1	1	1
Thread	2	2	2	2	2	2	2	2	2
Initialization time	0.000018	0.00002	0.000021	0.000022	0.000019	0.000018	0.000018	0.000024	0.000033
Thread compute time	3.523751	12.364939	3.080458	34.568468	10.484671	41.942321	53.135242	23.343529	10.070225
Write PNG time	1.152939	0.946612	0.765715	0.961856	1.481843	0.901876	1.539031	0.995231	0.96711
Total time	4.676708	13.311571	3.846194	35.530346	11.966533	42.844215	54.674291	24.338784	11.037368
Process	1	1	1	1	1	1	1	1	1
Thread	3	3	3	3	3	3	3	3	3
Initialization time	0.000019	0.000021	0.000019	0.00002	0.000022	0.000021	0.000018	0.000022	0.000019
Thread compute time	2.351429	8.252186	2.05586	23.07916	6.987645	27.976962	35.483878	15.578621	6.716372
Write PNG time	1.162455	0.947836	0.760895	0.968042	1.48168	0.900743	1.540511	1.001089	0.963097
Total time	3.513903	9.200043	2.816774	24.047222	8.469347	28.877726	37.024407	16.579732	7.679488
Process	1	1	1	1	1	1	1	1	1
Thread	4	4	4	4	4	4	4	4	4
Initialization time	0.000035	0.000021	0.000019	0.000022	0.000018	0.00002	0.000019	0.000022	0.000041
Thread compute time	1.766399	6.187515	1.551395	17.285039	5.232121	21.007626	26.625018	11.698381	5.036169
Write PNG time	1.157113	0.94754	0.753115	0.960356	1.483474	0.897758	1.532625	1.001179	0.958111
Total time	2.923547	7.135076	2.304529	18.245417	6.715613	21.905404	28.157662	12.699582	5.994321
Process	1	1	1	1	1	1	1	1	1
Thread	5	5	5	5	5	5	5	5	5
Initialization time	0.000035	0.000033	0.000041	0.00002	0.000031	0.000019	0.00002	0.000019	0.000021
Thread compute time	1.41486	4.95212	1.237482	13.837364	4.193321	16.804226	21.30085	9.357804	4.030731
Write PNG time	1.160397	0.942978	0.756135	0.956834	1.4834	0.898769	1.53933	1.002529	0.962724
Total time	2.575292	5.895131	1.993668	14.794218	5.676752	17.703014	22.8402	10.360152	4.993476
Process	2	2	2	2	2	2	2	2	2
Thread	1	1	1	1	1	1	1	1	1
Initialization time	0.00003	0.000034	0.000039	0.000025	0.000029	0.000029	0.00004	0.000027	0.000031
Thread compute time	7.028638	24.740941	6.137511	69.227421	20.953597	84.014062	106.520144	46.712085	20.129814
Write PNG time	1.16818	0.950249	0.762754	0.970395	1.48905	0.905065	1.546626	1.005642	0.968014
Total time	8.196848	25.691224	6.900304	70.197841	22.442676	84.919156	108.06681	47.717754	21.097859
Process	3	3	3	3	3	3	3	3	3
Thread	1	1	1	1	1	1	1	1	1
Initialization time	0.00005	0.000033	0.000033	0.00003	0.000027	0.000028	0.000036	0.000038	0.000028
Thread compute time	7.026272	24.736316	6.138182	69.184104	20.94536	84.023794	106.468744	46.785784	20.131487
Write PNG time	1.174938	0.954152	0.761314	0.967213	1.48566	0.903366	1.544881	1.014395	0.967731
Total time	8.20326	25.690501	6.899529	70.151347	22.431047	84.927188	108.013661	47.800217	21.099246
Process	4	4	4	4	4	4	4	4	4
Thread	1	1	1	1	1	1	1	1	1
Initialization time	0.000029	0.000028	0.00005	0.000029	0.000028	0.000031	0.000021	0.000027	0.000396
Thread compute time	7.026294	24.745348	6.143538	69.226752	20.943351	84.004105	106.502784	46.77548	20.117233
Write PNG time	1.1645	0.953958	0.762294	0.977723	1.482349	0.905031	1.568923	1.0155	0.965468
Total time	8.192823	25.699334	6.905882	70.204504	22.425728	84.910168	108.071708	47.791007	21.083097
Process	5	5	5	5	5	5	5	5	5
Thread	1	1	1	1	1	1	1	1	1
Initialization time	0.000028	0.000029	0.000034	0.000028	0.000036	0.000041	0.000049	0.000031	0.000023
Thread compute time	7.031562	24.755504	6.143788	69.20522	20.949939	84.010889	106.506712	46.810372	20.141182
Write PNG time	1.188505	0.958572	0.762239	0.971541	1.520258	0.912169	1.577946	1.017053	0.965183
Total time	8.220095	25.714105	6.906061	70.176789	22.470233	84.923099	108.084707	47.827456	21.106388

實驗簡單測了各個input在不同數量thread和process的情況下能加速多少

實驗結果也明確顯示multi-process在這種實作下是無法加速的(因為只有使用pthread)


上圖中multi-process的各個執行時間是取所有process中最慢的

其中**Initialization time**用來計算初始化區域變數的時間


Thread computation time用來計算從fork所有thread到join完所有thread所花時間

Write PNG time用來計算最後畫圖的時間


Thread	1	2	3	4	5
Initialization time	0.000019	0.000018	0.000019	0.000035	0.000035
Thread compute time	7.023348	3.523751	2.351429	1.766399	1.41486
Write PNG time	1.155882	1.152939	1.162455	1.157113	1.160397
Total time	8.179249	4.676708	3.513903	2.923547	2.575292
Speedup factor	1	1.74893301	2.327682067	2.797714215	3.176047221
Ideal speedup factor	1	2	3	4	5



Thread	1	2	3	4	5
Initialization time	0.00002	0.00002	0.000021	0.000021	0.000033
Thread compute time	24.733721	12.364939	8.252186	6.187515	4.95212
Write PNG time	0.947597	0.946612	0.947836	0.94754	0.942978
Total time	25.681338	13.311571	9.200043	7.135076	5.895131
Speedup factor	1	1.929249222	2.791436736	3.599308262	4.356364261
Ideal speedup factor	1	2	3	4	5





Thread	1	2	3	4	5
Initialization time	0.000018	0.000021	0.000019	0.000019	0.000041
Thread compute time	6.139337	3.080458	2.05586	1.551395	1.237492
Write PNG time	0.762026	0.765715	0.760895	0.753115	0.756135
Total time	6.901381	3.846194	2.816774	2.304529	1.993668
Speedup factor	1	1.794340327	2.450101073	2.994703473	3.461650084
Ideal speedup factor	1	2	3	4	5





Thread	1	2	3	4	5
Initialization time	0.00002	0.000022	0.00002	0.000022	0.00002
Thread compute time	69.2553	34.568468	23.07916	17.285039	13.837364
Write PNG time	0.968045	0.961856	0.968042	0.960356	0.956834
Total time	70.223365	35.530346	24.047222	18.245417	14.794218
Speedup factor	1	1.976433469	2.92022775	3.848822145	4.74667637
Ideal speedup factor	1	2	3	4	5

strict31.txt



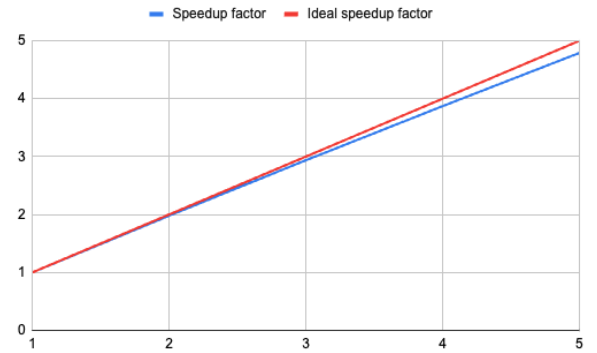

Thread	1	2	3	4	5
Initialization time	0.000022	0.000019	0.000022	0.000018	0.000031
Thread compute time	20.96561	10.484671	6.987645	5.232121	4.193321
Write PNG time	1.48397	1.481843	1.48168	1.483474	1.4834
Total time	22.449602	11.966533	8.469347	6.715613	5.676752
Speedup factor	1	1.876032264	2.650688654	3.342896918	3.954656113
Ideal speedup factor	1	2	3	4	5

strict32.txt




Thread	1	2	3	4	5
Initialization time	0.00002	0.000018	0.000021	0.00002	0.000019
Thread compute time	83.905639	41.942321	27.976962	21.007626	16.804226
Write PNG time	0.900907	0.901876	0.900743	0.897758	0.898769
Total time	84.806566	42.844215	28.877726	21.905404	17.703014
Speedup factor	1	1.979416964	2.936746682	3.871490615	4.790515672
Ideal speedup factor	1	2	3	4	5


strict33.txt




Thread	1	2	3	4	5
Initialization time	0.000019	0.000018	0.000018	0.000019	0.00002
Thread compute time	106.299448	53.135242	35.483878	26.625018	21.30085
Write PNG time	1.541206	1.539031	1.540511	1.532625	1.53933
Total time	107.840673	54.674291	37.024407	28.157662	22.8402
Speedup factor	1	1.972420145	2.912691431	3.829887332	4.721529277
Ideal speedup factor	1	2	3	4	5



Thread	1	2	3	4	5
Initialization time	0.000017	0.000024	0.000022	0.000022	0.000019
Thread compute time	46.753523	23.343529	15.578621	11.698381	9.357604
Write PNG time	1.000165	0.995231	1.001089	1.001179	1.002529
Total time	47.753705	24.338784	16.579732	12.699582	10.360152
Speedup factor	1	1.96204153	2.880245893	3.760258015	4.609363357
Ideal speedup factor	1	2	3	4	5



Thread	1	2	3	4	5
Initialization time	0.00002	0.000033	0.000019	0.000041	0.000021
Thread compute time	20.139368	10.070225	6.716372	5.036169	4.030731
Write PNG time	0.95577	0.96711	0.963097	0.958111	0.962724
Total time	21.095158	11.037368	7.679488	5.994321	4.993476
Speedup factor	1	1.91124895	2.746948494	3.519190581	4.224543785
Ideal speedup factor	1	2	3	4	5



以上幾張圖顯示初始化時間根本可以忽略。當每個thread的computation time很大時，write PNG time的佔比會非常小。

因此可以發現當總執行時間非常長時，我實作的平行化計算能夠逼近ideal speedup factor

也就是說如果撇除掉initialization time和write PNG time這兩個我沒做平行化的部分，其餘時間的平行化程度很高(可以從strict34.txt看出來)

hybrid

經過pthread的實驗後，發現只要測strict34.txt就可以展現這隻程式的平行度優化結果(因為可平行化的部分佔比很大)

Initialization time用來計算初始化區域變數的時間


OpenMP time用來計算#pragma內的程式執行時間，由於這段會同時受process和thread數量影響，所以直接把這段當成是可平行化的最大部分

Communication time用來計算MPI_Gatherv所花時間

Gather image time用來計算把MPI_Gatherv組起來的圖片重組花的時間

Write PNG time用來計算最後畫圖的時間

Node	1	1	1	1	1	1	1	1	1	1	1	1	1
Process	1	2	3	4	5	6	7	8	9	10	11	12	
Thread	1	1	1	1	1	1	1	1	1	1	1	1	1
Initialization time	0.572333	0.742797	0.754974	0.656889	0.658948	0.778726	0.77498	0.690338	0.670201	0.682432	0.803367	0.812887	
OpenMP time	104.159533	52.178636	34.896724	26.07274	20.862187	17.381102	14.893995	13.059597	11.589486	10.433455	9.488143	8.693935	
Communication time	0.026706	0.02869	0.029594	0.036985	0.027696	0.030364	0.030447	0.029455	0.047198	0.031879	0.042589	0.033835	
Gather image time	0.032135	0.027505	0.026384	0.029524	0.027305	0.032159	0.031954	0.026717	0.028126	0.027432	0.028417	0.030039	
Write PNG time	1.552002	1.548266	1.537825	1.55069	1.547066	1.5514	1.539039	1.550449	1.55265	1.549093	1.549196	1.553284	
Total time	106.342709	54.525894	37.245501	28.346828	23.123202	19.773751	17.270415	15.356556	13.887661	12.724291	11.911712	11.12398	
Speedup factor	1	1.950315734	2.85518267	3.751485316	4.598961208	5.377973506	6.157507448	6.924906144	7.657352019	8.357456537	8.927575566	9.559771682	
Ideal speedup factor	1	2	3	4	5	6	7	8	9	10	11	12	
Speedup computation	1	1.996210346	2.984794017	3.994959218	4.992742755	5.992688668	6.993391162	7.975708056	8.987416094	9.983225403	10.9778629	11.98071219	




上圖是固定1個thread，實驗process從1~12個

左圖是上述各種耗時的佔比圖

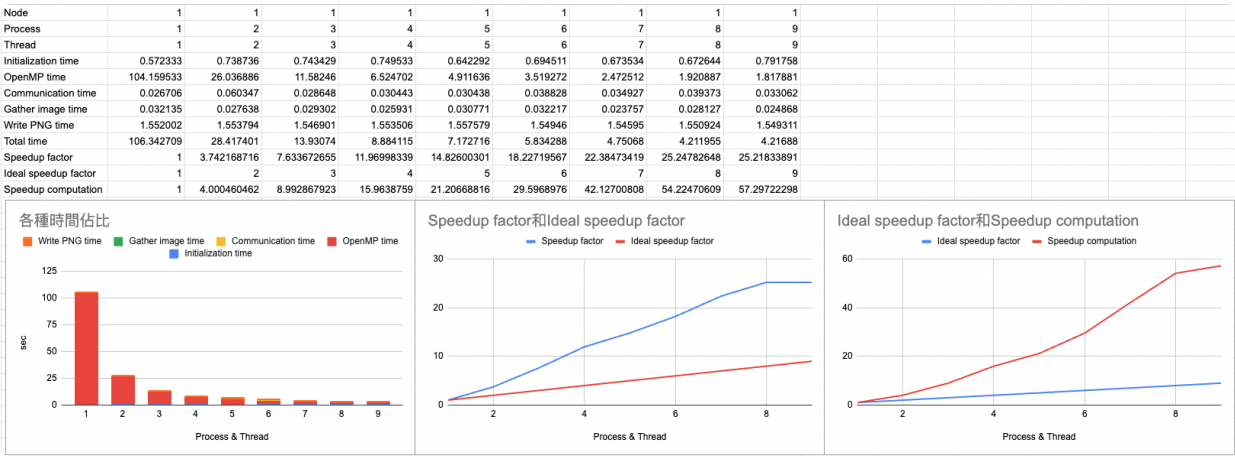
中間是總時間的speedup factor和ideal speedup factor的比較

右圖是如果不算無法平行的部分，也就是只算#pragma內的程式的speedup是否可以逼近ideal speedup factor

Node	1	1	1	1	1	1	1	1	1	1	1	1	1
Process	1	1	1	1	1	1	1	1	1	1	1	1	1
Thread	1	2	3	4	5	6	7	8	9	10	11	12	
Initialization time	0.572333	0.573443	0.57406	0.573721	0.572041	0.56985	0.58726	0.577853	0.581018	0.583119	0.571049	0.57586	
OpenMP time	104.159533	52.740087	34.715778	26.065286	20.967251	17.409615	15.012787	13.030434	11.58508	10.607236	9.639782	8.698177	
Communication time	0.026706	0.028174	0.033951	0.026475	0.026754	0.032457	0.026429	0.031341	0.030468	0.026405	0.031744	0.032265	
Gather image time	0.032135	0.025645	0.032495	0.031784	0.031799	0.028885	0.031998	0.031327	0.03176	0.026511	0.031411	0.028462	
Write PNG time	1.552002	1.55098	1.539508	1.549659	1.567844	1.551151	1.555772	1.550404	1.554634	1.555677	1.541575	1.550388	
Total time	106.342709	54.922447	36.895792	28.246925	23.165689	19.591958	17.214246	15.221359	13.78296	12.798948	11.815561	10.885152	
Speedup factor	1	1.936233996	2.882244918	3.764753473	4.590526489	5.427875509	6.177599007	6.986413565	7.715520396	8.308707012	9.000225127	9.769519893	
Ideal speedup factor	1	2	3	4	5	6	7	8	9	10	11	12	
Speedup computation	1	1.974959446	3.000351396	3.996101673	4.96772481	5.982874004	6.938054407	7.993558235	8.990834159	9.819667725	10.80517516	11.97486933	



上圖是固定1個process，實驗thread從1~12個



上圖是同時增加process和thread的數量看看speedup factor是否能呈指數成長
實驗結果大概在5個process和5個thread時就上不去了

Profile

pthread



從profiler可以看出幾個點：

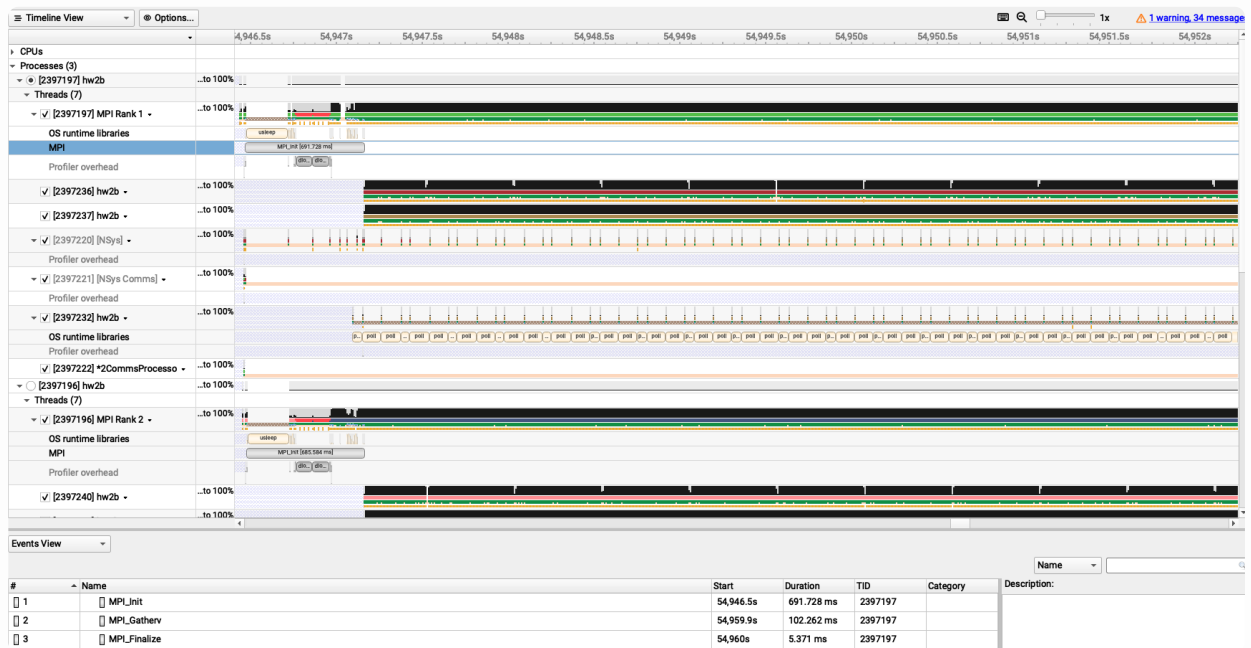
- 1. 有些thread中間會有一小段空白(約8ms)，這些應該是做完工作準備拿新工作時，別的thread已經搶到lock，所以是等待lock的時間。
- 2. 有時候thread在等lock，原本在使用的CPU core會被其他thread搶佔，所以等拿到lock後，會交換執行的CPU core。

Events View				
#	Name	Start	Duration	TID
6	pthread_join	0.0305221s	21.387 s	1974503
4372	fclose	22.9377s	24.129 ms	1974503
8	open	21.4182s	5.495 ms	1974503
7	pthread_join	21.4177s	484.412 μs	1974503
4	pthread_create	0.0302433s	169.291 μs	1974503
1	pthread_create	0.0297675s	158.784 μs	1974503
2	pthread_create	0.0299278s	139.272 μs	1974503
3	pthread_create	0.0300952s	134.617 μs	1974503
5	pthread_create	0.0304147s	101.582 μs	1974503
2429	fwrite	22.2721s	12.026 μs	1974503
2114	fwrite	22.1634s	11.742 μs	1974503
2146	fwrite	22.1739s	11.699 μs	1974503
3217	fwrite	22.5435s	11.250 μs	1974503
350	fwrite	21.5423s	11.069 μs	1974503
10	fwrite	21.4248s	11.062 μs	1974503
1705	fwrite	22.019s	10.906 μs	1974503
791	fwrite	21.6957s	10.905 μs	1974503
2902	fwrite	22.4352s	10.838 μs	1974503
1642	fwrite	21.9966s	10.824 μs	1974503
476	fwrite	21.5862s	10.750 μs	1974503
2744	fwrite	22.3801s	10.726 μs	1974503
917	fwrite	21.7393s	10.706 μs	1974503
4131	fwrite	22.856s	10.690 μs	1974503
382	fwrite	21.5536s	10.679 μs	1974503
1012	fwrite	21.7721s	10.678 μs	1974503
2587	fwrite	22.3258s	10.652 μs	1974503

Events View			
#	Name	Start	TID
1	[kernel.kallsyms]0xffffffff758cb00	0.0305405s	1974518
2	hw2alcalculate(...)	0.030976s	1974518
3	hw2alcalculate(...)	0.0314115s	1974518
4	hw2alcalculate(...)	0.0318469s	1974518
5	hw2alcalculate(...)	0.0322822s	1974518
6	hw2alcalculate(...)	0.0327175s	1974518
7	hw2alcalculate(...)	0.0331528s	1974518
8	hw2alcalculate(...)	0.033588s	1974518
9	hw2alcalculate(...)	0.0340233s	1974518
10	hw2alcalculate(...)	0.0344585s	1974518
11	hw2alcalculate(...)	0.0348937s	1974518
12	hw2alcalculate(...)	0.0353291s	1974518
13	hw2alcalculate(...)	0.0357643s	1974518
14	hw2alcalculate(...)	0.0361995s	1974518
15	hw2alcalculate(...)	0.0366349s	1974518
16	hw2alcalculate(...)	0.0370701s	1974518
17	hw2alcalculate(...)	0.0375053s	1974518
18	hw2alcalculate(...)	0.0379405s	1974518
19	hw2alcalculate(...)	0.0383758s	1974518
20	hw2alcalculate(...)	0.0388111s	1974518
21	hw2alcalculate(...)	0.0392466s	1974518
22	hw2alcalculate(...)	0.0396819s	1974518
23	hw2alcalculate(...)	0.0401174s	1974518
24	hw2alcalculate(...)	0.0405528s	1974518
25	hw2alcalculate(...)	0.0409882s	1974518
26	hw2alcalculate(...)	0.0414236s	1974518
27	hw2alcalculate(...)	0.041859s	1974518
28	hw2alcalculate(...)	0.0422944s	1974518
29	hw2alcalculate(...)	0.0427297s	1974518
30	hw2alcalculate(...)	0.0556618s	1974518
31	hw2alcalculate(...)	0.0560969s	1974518
32	hw2alcalculate(...)	0.0565321s	1974518
33	hw2alcalculate(...)	0.0569674s	1974518

每個thread做完一次工作大概是0.0004352秒

hybrid




Hybrid方法就不像pthread會用lock去管理工作，而是在MPI_Init後就針對不同rank的process分配好需要做的工作，因此不會有中間的小斷點

Discussion

Scalability

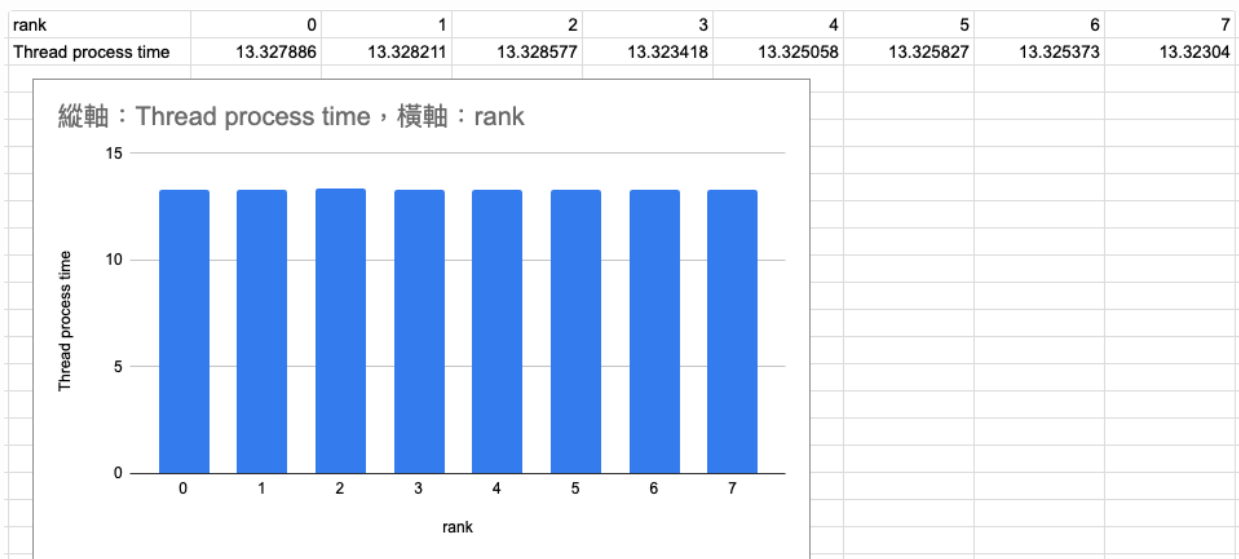
拿strict34.txt當作input來計算從1個thread scale up到12個thread。



實驗結果發現當thread數量逐漸scale up，speedup factor逐漸趨緩，看起來會收斂到一個最終值
 但是我又另外做了一個實驗，只計算可平行化的部分speedup會不會也收斂
 結果發現當process在12個以下時，沒辦法看出明顯得收斂，甚至接近ideal speedup factor
 因此可以得到一個結論，撇除可平行化的部分，剩餘不可平行化的部分就是無法逼近ideal speedup factor的bottleneck

Load balancing


拿strict34.txt當作input來實驗8個thread的執行時間



Pthread實驗結果顯示，每個thread執行時間很平均，算是達成一個不錯的load balance

藉由控制blockPerThread來達到更fine-grained的工作分配是可行的

然而我做了許多小實驗，從blockPerThread = 30，做到blockPerThread = 2500，跑judge的總時間是漸減的。由此可知，當工作被切得越小份時，雖然能增加計算資源的利用，但是每次重新分配資源到vector所產生的overhead也會相對提升。因此當圖片越大時，blockPerThread就要相對上升，以減少overhead的產生



Hybrid方法因為rank0要負責做圖片的aggregate，所以會花比較多時間，因此會成為效能瓶頸
其餘的process執行時間都差不多，算是達成不錯的load balance
以目前簡單做的小實驗來看，data parallelism的schedule用static、branch = 10效果最好

Conclusion

經過這次作業，我了解到不是每個程式都適合拿來做平行計算優化。這次作業的程式就很適合用平行計算優化。從sequential的程式來看，每個計算工作都互相獨立，且性質相同，非常適合同時做計算。不論是每個工作的計算能平行，還是最後畫圖的部分，我相信都是可以用平行計算去逼近ideal speedup的。

實作上遇到比較大的困難大概是看不到vector裡的值，所以debug的時候會需要把vector裡的所有元素都倒出來，會需要看一堆log來檢查每個工作是不是都做了，也需要檢查每個工作是不是在正確的時間結束工作。